

MIT 6.858 - Computer Systems Security

INTRODUCTION, THREAT MODELS

- Security
 - Policy - what you want to achieve (confidentiality, integrity)
 - Threat Model - assumptions about bad guy
 - Mechanism - sw/hw implementation
- Negative goal - e.g. nobody but X can access Y
- Adding security questions to passwords weakens security, more vulnerabilities
- Threat Model weaknesses
 - Humans
 - Time - computers constantly improving
- SSL Certificates
 - Written in C, null terminated strings
 - can request google.com\0x.foo.cum\0, stops at first \0
- Buffer overflow
 - stack grow down - overflow w/ "outer" method (caller of this)
 - stack grow up - overflow w/ "inner" method (called inside method)

2. CONTROL HIJACKING ATTACKS

- Buffer overflows
 - Take advantage of
 - System software is written in C
 - C exposes raw memory addresses
 - No bounds checking
 - Knowledge of x86 (stack, registers, calling conv.)
 - Can run with priority of the function it overtakes
 - Hardware is watching you all the time, not OS
- Avoiding Buffer Overflow
 - Avoid bugs in the first place
 - Don't use gets(), gotos, etc. → although these are not always used
 - Build tools to fix bugs
 - Static analysis
 - Fuzzing
 - Use memory safe languages
 - Usually not possible because
 - Legacy
 - Need low level access
 - Performance (less of an issue now)

- Buffer overflows exploit
 - Gaining control instruction pointer
 - Make it point to malicious code
- Stack canaries
 - Put a value before return address (canary), check it before jumping to return address
 - Problem - attacker can guess/find canary
 - One solution is to have canary be /0, carriage return, line feed,
 - Another solution is a randomized value
 - When they fail
 - Overwriting pointer functions

```

int *ptr = m; ← overwrite here
...
*ptr = 5;
  
```

- If you can guess a random canary
- malloc and free attack
 - overflow overwrites size data of malloc block

- Bounds checking
 - Hard to tell what you want, esp with structs + unions
 - For a pointer p' that's derived from p , then p' 's should only be used to deref. memory that belongs to p

- Electric fences
 - set guard page on heap, if pointer hits guard page if fail
 - very space intensive because it takes a whole page

- Fat Pointers
 - Add bounds info to pointer
 - 3x size: start, end, current
 - Can cause concurrency bugs

- Shadow data structure
 - For each object, store size too
 - Can't check mathematical bounds by itself because for loops, etc must be broken once and become invalid to stop

- Buddy Allocation
 - Treats unallocated memory as one big block
 - Splits by powers of two until memory is just big enough
 - Splits until actual size $\geq \frac{1}{2}$ size of block

28	20	50	
0	32	64	128

→ when freed, combines 2 of same size
WASTEFUL!

Baggy Bounds

- Round up each allocation to power of 2, align request to power
- Express each bound as \log_2 of allocation size
- Store limit info in linear array with 1 byte/entry
- Allocate mem. at slot granularity

WITH SLOT SIZE 16 FOR EXAMPLE

$$p' = p + i;$$

$$\text{size} = 1 \ll [p \gg \log \text{of slot size}]$$

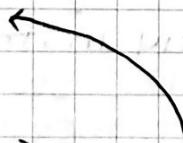
$$\text{base} = p \& \sim(\text{size} - 1) \leftarrow \text{clear bits that are offset of base}$$

- USES VIRTUAL memory system to prevent OOB dereference

3. BUFFER OVERFLOW EXPLOITS AND DEFENSES

- Baggy bounds throws error if you get beyond half the slot
 - If out of bounds but < half, set high order bit so if it's dereferenced it throws a hard error
- When we access the table the entry # means pointer size is 2 raised to that number
- Max is automatically set to 31
- Instrumented code
- Can't detect a bad pointer from instrumented to uninstrumented
- Regular pointer

zero	size	addr
21	5	38



} still 64bit addresses

• OOB

offset	size	zero	addr
13	5	8	38

• Still problems

- Can't look at uninstrumented code
- Could still be within baggy bounds if pointer

• Costs of baggy bounds

- Space - store bounds table
- CPU overhead to check bounds
- False alarm - make OOB ptr, but doesn't dereference
- Need compiler support

- Non-executable memory
 - Paging hardware writes 3 bits for each page (R, W, X)
 - W^X (can't do both)
 - Makes it hard to dynamically write code
 - Just in time compiler - turn common func into x86
- Randomized address spaces
 - A lot of attacks use hard code addresses
 - Stack randomization
 - they can just guess seed
- Heap Attack
 - Stuffs a bunch of shell code into memory
 - Jump to random
 - No-op sleds
 - Tons of no-ops, actual code at end so you just fall through to there
- Change system call numbers to be random (e.g. Stop is 4, etc)
 - Some people suggest a hardware change, adding a XOR key that's XOR'd with each instruction
- GCC, VScode use canaries
 - Linux, Windows use non exec memory
- Return oriented programming
 - You can fake a calling frame, put it in the buffer
 - Gadgets
 - Small set of assembly instructions
 - Works with non-executable stack
 - Uses stack pointer as instruction pointer
- Guessing canaries
 - Assume server has BOF vulnerability
 - Server will crash and restart on bad canary
 - After restart, canary + addresses are not reset
 - Fork inherits canary, etc
 - Probe each next byte
- Defeat by catching segfault and staying open
 - Attacker thinks guess is correct
- 32 bit machine, arguments are in stack
 - 64 bit - registers

• Blind Return Oriented Programming

- Find stop gadget - pause program but not crash
- Find gadgets that pop stack entries
- Example

STOP : 0x5

TRAP : 0x0

PROBE: 0x4..8 ← say this is pop rax, ret

! If probe works, pops trap, returns to stop, all good. If not, returns to trap, crashes ↑

To find more pops, just add more traps in between

- Know a pop happens, but not which reg its popped to
- Take advantage of 'pause' call, put it in between gadgets, get it to hang and inspect registers
- Try to find a gadget that pops pause # to rax, then call syscall(), if it's right # it will pause, else crash
- Invoke Wintel)
 - Need pop rdi, rsi, rdx, rax, invoke syscall()
 - ↓ ↓ ↓ ↑
 - socket buf buf len syscall #

• How to defeat ROP

- Use exec() instead of fork() so values are re-randomized
- Use windows - no fork() equivalent
- When crash happens, don't close connection
- NOT A GOOD IDEA - hashing time of day
 - Bits of entropy are reduced if input is guessable (e.g. time)

4. PRIVILEGE SEPARATION

- If app that attacker gains access to has privileged status, it can affect much more besides itself
 - Privilege separation aims to break app into smaller pieces, only give them access to what they need
 - VMs are used for this
 - Unix provides you with tools

• UNIX

- Principal: What privileges the subject has [userid]
[groupid]
- Subject: process - single uid, list of gid
- Object: What to protect (files, dirs, sockets, other proc)
 - Files - R, W, X, chmod
 - Directory - unlink, link, rename, create, lookup {inode
 - owned by uid/gid - can change perm if uid is same
 - permission bits for uid, gid, other

open ("/etc/passwd") - need x /, x /etc, r /etc/passwd

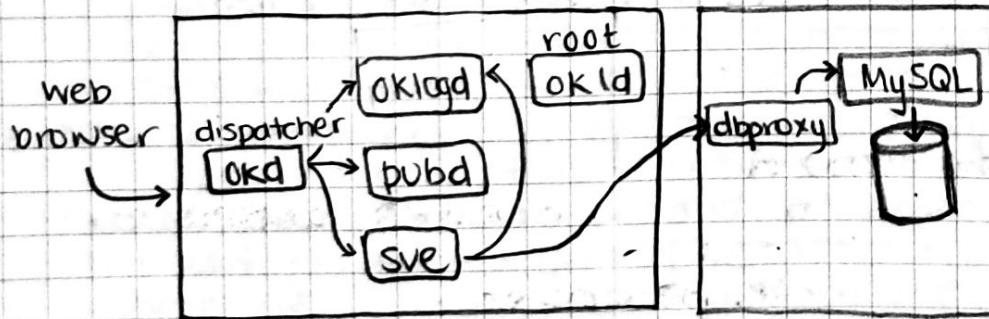
- can combine permissions in layers

e.g. E 6.858 gid, T A gid, ≠ 6.858 TA gid

Make /foo/bar/file, give foo 6.858 perm, bar TA perm

- File descriptors - allows you to open file
 - Can be passed to you
- Process - create, kill, debug (ptrace)
 - Needs same UID
- Networking - listen, connect, read/write, send/receive
 - Anyone can connect
 - Need to be SU to listen to port < 1024 (uid = 0)
- Creating UIDs
 - Can set-uid if root (uid=0), also set-gid and set-groups
 - Login process runs as root, checks password from /etc/passwd and /etc/shadow
- Set uid binaries
 - When you exec, switches uid of process to owner of binary
 - Have to be careful with what environment variables you pass
- How to avoid setuid issues
 - Call chroot("/foo")
 - / = /foo (restrict namespace)
 - Prevents ... bc its now /foo/.../
 - Need to be uid 0

• OKWS: Case Study



- Need token to call dbproxy
- Each process has a specific uid

- OKd and Sve share same chroot
- OKld needs root bc port 80
- Why don't they have true userids instead of service ids?
 - Performance overheads
 - Depends on having access to all data at once

6. CAPABILITIES

- capabilities let you control privileges better than Unix
- 2 examples: confused deputy, sandboxing

- Confused deputy
 - Had FORTRAN compiler in /sysx/fort
 - Wanted it to write usage statistics into /sysx/stat
 - Home files license - gives you extra privileges, can access files in /sysx/*
 - Someone told a program to output to /sysx/bill, for example, since the compiler had Home files License even though user didn't
 - 2 sources of privilege
 - User Invoking } their design used a union of this
 - Home files license }
 - What goes wrong?
 - Ambient authority - making an operation, but decision of whether it succeeds comes from outside parameters
 - Access control checks are complicated
 - If you try to replicate check system, maybe you miss some, or kernel updates and adds new checks and your app doesn't keep up
 - Capability has no ambient authority
 - Capability points to object, if you have capability you can access object

- Capability
 - Need to be unforgeable
 - File descriptor - just an integer
 - Kernel interprets based on descriptor table
 - To forge a file descriptor, you have to pick a number that refers to a non-null entry in table, but those are already the capabilities you have
- Solving Confused Deputy
 - Don't pass file names around
 - Frontend turns input files into file descriptors, passes them
 - If a user can't open it, it fails
 - Interpreter then deals with file descriptors
- Capabilities - can't talk about object without conveying rights to that object
- Privileges don't stack

- Sandboxing
 - Network stuff - how much could a hacker control
 - Best if dev knows it will be run in sandbox
 - Why not use VM?
 - Memory overhead
 - Network issues
 - Hard to share data

- How to sandbox
 - Unix
 - Change permissions, make sure there aren't world readable files
 - capsicum
 - Once you enter sandbox, everything depends on capabilities
 - Global namespace (e.g. file system)
 - Can't make it just a library, needs to be a kernel thing
 - Could still make syscall
 - When you enter capsicum, stops all syscalls to global namespaces
 - Can't open sockets, need to set them up beforehand
 - How do you kill processes?
 - Using a file descriptor
 - What about using .. ?
 - Disallow, just to be safe
 - Once you enter capability mode, does UID matter?
 - Yes to be compatible w/ general Unix permissions
 - Why do they need a library?
 - To use function lch - start (vs. cap - enter)
 - Erases all inheritance stuff, start a process fresh
 - Fd lists
 - Generalization of how Linux handles file descriptors
 - Don't allow setuid binaries
 - In general, can't gain any privileges

7: SANDBOXING NATIVE CODE

- Native Client
 - Software fault isolation
- Why do you need to run native code? (In the case of browser)
 - Performance
 - Existing "Legacy" code
 - Don't need to use Javascript
 - Other languages
- How to use Native Client
 - Need some JS to connect webpage to Na. Cl. module
 - Can contain crashes, module goes down but page doesn't
- Alternatives
 - Trust the developer/ask user
 - Insecure, can't tell if binary is safe (as user)
 - Native client ensures it's safe, and protects you
 - OS/Hardware Isolation
 - Worried about OS bugs

- OSs could be incompatible, need diff isolation mechanism
- Hardware bugs

• Software Fault Isolation

- Don't rely on OS/HW to see if it's safe at runtime, look ahead to see if it's safe
 - Safe - allow (e.g. ALU, math, movs)
 - Unsafe - (e.g. mem access, syscall, privileged instruction)
 - Necessary - (e.g. mem access) instrument, add in checks, make safe
 - Unnecessary - prohibit, kill application
- Since it can't access memory, network, etc., need to add a Trusted Service runtime
 - Not checked
 - things like alloc, threads

• Safety

- No disallowed instructions can execute
 - syscall, privileged instructions
- All code/data accesses are in bounds
 - [0... 256 MB] in module is safe
 - without, you could do disallowed instructions by jumping wherever
- This means one module per process since you need low addresses

• How to Implement

- Scanning instructions
 - Problem: variable length instructions (need to look at the first couple of bytes to see how long it will be)
 - Could also jump to middle of instruction
 - Look at every offset?
 - No bc false positives
- How they actually do it
 - Need to deal with jumps
 - So they check where the jumps go
 - If it jumps somewhere we haven't seen (e.g. we saw 0x20 but not 0x22 bc instruction is 4 bytes long not 2) flag it
 - Issue - indirect jumps (e.g. jmp %eax)
 - Solve with instrumentation
 - Add an 'AND 0xfffffe0' to clear lowest 5 bits, forces value to be a multiple of 32 (other threads can't break)
 - Then make sure everything at addresses that are multiples of 32 are safe
 - Why 32? Power of 2, fits longest instruction, but not too big
 - Prevents jumping over and by treating the and+jmp as one instruction

- Rules

- C1 - binary is not writeable after loading
- C2 - linked @ 0, starts @ 64K
- C3 - indirect jumps use and + jmp
- C4 - pad out to page boundary w/ halt (safety buffer)
- C5 - no instructions span over 32 byte boundary
- C6 - all instr. reachable by disassembly
- C7 - all direct jumps are OK

- Bounds checking

- zero out upper bytes

- Segmentation in x86

- Whenever a process is running, hardware makes segment descriptor table

- segments numbered

- each entry has base and length

- whenever x86 talks about memory, it uses this table, talking about a specific segment

- Segment Selectors (16 bit) are cs, ds, es, fs, gs, ss

- cs - code selector

- ds/es - for memory

- ss - stack

- Manipulate table with system calls

- Tables are per process

- How to jump out of Native Client

- change the selectors

- Trampolines and springboards

- Trampoline code comes from trusted runtime, not NaCl

- Trampoline

- move values into selectors

- jump into normal runtime

- Springboard - jump back into NaCl

- restore selectors

- jmp into NaCl module

- put halt at front so NaCl can't jump into the springboard itself

8: WEB SECURITY MODEL

- Browser is complicated today

- JavaScript

- DOM Model

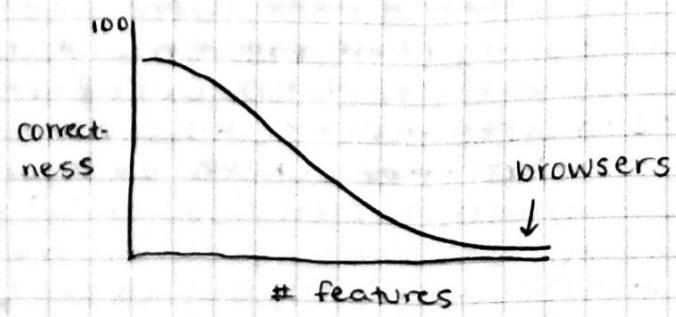
- XMLHttpRequest (AJAX)

- Web Sockets

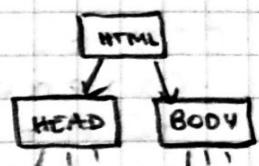
- Multimedia support

- Geolocation

- NativeClient



- Parsing contexts
 - <script> var x = "UNTRUSTED"; </script>
 - Could put in JS
 - Could put in <script> and add in HTML
- Same-origin policy
 - 2 websites should not be able to tamper with each other unless they want to
 - strategy
 - Each resource is assigned an origin
 - JS can only access resources from its own origin
 - Origin: network protocol scheme + hostname + port
 - http://foo.com/index.html = HTTP, foo.com, 80
 - https://foo.com/index.html = HTTPS, foo.com, 443
 - Like a UID
 - 4 ideas
 - Each origin has client-side resources
 - Cookies - tiny file associated with origin
 - DOM Storage - key/value storage
 - JS namespace
 - DOM tree - tree of HTML structure
 - Visual display area
 - Each frame gets origin of its URL
 - Scripts execute with the authority of its frame's origin
 - Passive content gets 0 authority
 - Non executable things like Images/CSS
 - postMessage() allows 2 frames to communicate



DOM tree

- MIME Sniffing Attack
 - MIME type - text/html, jpg, etc.
 - Web servers could misattribute MIME types
 - IE used to look at first 256 bytes to guess what type of object it is
 - Could coerce an image into HTML/JS content + execute it
- Frames and Window Objects - Same Origin Policy
 - Frame: instance of a DOM node
 - Window: alias for global namespace
 - contain pointers to one another
 - Get origin of frame's URL OR get suffix of original domain
 - Two frames can access each other if
 - Both frames set document.domain to same val
 - Neither have changed document.domain + vals match
 - Prevents domains from being attacked by subdomains
- Cookies
 - Have domain + path

- Can be set by server or on client side
 - Client can change cookie w/ js document.cookie
- Frames accessing other cookies
 - If they have same doc.domain and both changed it
- XMLHttpRequest
 - Can only generate one if its going to its origin server
 - Same origin unless Cross Origin Request (CORS) enabled
 - Adds HTTP header called access-control-allow-origin:<URL>
- Images
 - Frames can load images from wherever
 - Can't directly inspect bits
 - You can still get, say, size info based on what it affects
- CSS
 - Can embed from any origin
 - Can't inspect file
 - can infer based on what nodes are created
- JavaScript
 - can cross origin fetch
 - Can't inspect source code
 - Since functions are objects though, you can call toString
 - People try to hide by obfuscating and min-ifing
- Plugins
 - Almost out of use
- Cross Site Request Forgery (CSRF)
 - Example: http://bank.com/transfer?amount=5000&to=hacker
 - Has all the users cookies, so it'll run
 - Fix by introducing randomness the attacker can't guess
 - Add a random token as a query
- Network Addresses
 - DNS rebinding attack
 - Run malicious JS with authority of the website
 - Steps
 - Register a domain name (e.g. attacker.com)
 - User visits attacker.com
 - Browser generates DNS request to attacker.com, attacker responds with something that will expire soon (needs to be re-checked soon)
 - Attacker binds attacker.com to user's IP address
 - Website fetches new object, but this goes to victim.com
 - Can continually rebinding

- Fixes
 - Make it impossible to resolve external domains to internal IPs
 - Could also ignore TTL and keep for, say, 30 min

- Clickjacking Attack

- Pixels are bounded by frames
- Parent frames can write over child frames
- Hide malicious frames behind normal ones

- Solution

- Frame busting code
 - if (self != top) ← JS
- X-Frame-Options HTTP header
 - Don't allow anyone to put my content inside a frame

- Typosquatting Attacks

- Hacker can buy Cats.com with Cyrillic C instead of Latin C

- Plugins

- Java assumes same IP, diff host = same origin

- Screensharing

- Can screenshot outside of the frame

9: SECURING WEB APPLICATIONS

- Shellshock

- You can specify custom HTTP headers
- (){:;};

- XSS

- Malformed URL can bypass XSS protection.

- Defense

- XSS filters in the browser
 - Checks if <script> is in URL
 - Can't prevent persistent XSS

- HTTP Only Cookies

- Privilege separation
 - Use a separate domain for untrusted data

- Content Sanitization

- Don't trust any user input

- < → <

- " → "

- Content Security Policy

- Tell browser what content can be loaded and from where

- Sent in HTTP Header

• SQL

- Escape SQL queries
- Solution - rigorously encode your data
 - Character escaping

• COOKIES

- Session management
- Stateless cookie - avoid attacks?
 - Need to then authenticate every request
 - MAC - Message Authentication Codes
 - Hash key with message
 - USE DOM client side storage
 - LESS SUSCEPTIBLE TO SOP ATTACKS
 - Client Side Certificates
 - Revocation is hard

• Protocol Vulnerabilities in Web Stack

- URL Parsing
 - `http://example.com:80@foo.com`
 - Which is the origin?
 - Flash says example
 - Browser says

• JAR - GIFAR attack

- Can combine .gif and .zip, bc .gif reads top down, - zip reads bottom up
- JAR files are derivatives of .zip
- Embed JS in GIF
 - Can be used in conjunction with XSS attack
- Also works with PDF

• Covert Channel Attack

• CSS-based Sniffing Attack

- Attacker wants to know what websites user visited
- Get them to visit a site the attacker controls
- Generate a list of URLs the user may have visited, then use JS to see what color they are
- Solution - Browser lies to JS about link colors

• Cache-Based Attack

- Time response time for various websites
- Shorter time = cached = visited
- can use this on something like Google Map tiles
 - Visits = maybe you live there

• DNS based attack

- Check how long DNS requests take
- Shorter time = you visited
- Caching doesn't matter since DNS is OS stuff

• Rendering attacks

- Faster to render a URL you have visited
- When an iframe is loading, attacker owns it
- When they lose it, they know it rendered

10: SYMBOLIC EXECUTION

$$t_0 = \begin{cases} x > y \Rightarrow x \\ x \leq y \Rightarrow y \end{cases} \quad \text{when } t_0 < x$$

can also add preconditions, like $x \leq y$

SMT - will either return satisfying conditions or unsat.
↳ Satisfiability Modulo Theories

11: VR/WEB

- Full stack web development language

12: NETWORK SECURITY

• Threat Model

- Intercept, modify packets
- Inject packets
- Attacker has computer
- Distributed

• TCP

- 3 Way Handshake 32 BIT
- Client: C → S SYN(SEQUENCE NUM_c)
- Server: S → C SYN(SN_s), ACK(SN_c)
- Client: C → S ACK(SN_c)
- SN is incremented per byte sent

- The SN was easy to guess

• The starting SN is bumped up by some 25, 64, 128K

• Attack: Forge another machines' IP

- IP Based Authorization

- SMTP still uses this
- rlogin used this

• RST

- If you know a sequence number, you can send a RESET request

- Multiple routers connecting to each other

• If a TCP connection breaks, reset routing tables

• DOS Attack

• Routers ensure it's from the immediately next router by only accepting packets with TTL of 255 (max)

• So you can see who's sending reset packets

• Data Injection

- Inject 32-bit offset existing connection

• Prevention

- Generate 32-bit offset to add to initial sequence number

• DNS

- Uses UDP - stateless

- Just need SRC + DST IP and port (53)

- Inject packets before you get your real response

- It might even cache this as the correct IP

• Solution

- Add 16-bit query ID

- Add 16-bit random + 53

• Non Existence Records

- Have to be signed by server

- NSEC - organizes domains alphabetically, if there's nothing in between, you know it doesn't exist

- You can query for all the "gaps" to get the things on either side

- e.g. [aaa - bbb] query for aab, doesn't exist, tells you what's on left and right (aaa, bbb). Then query for bbc, tells you bbb and ccc, continue ad nauseum

• SYN Flooding

- Data structure with connection, SN_c and SN_s

- DOS by sending a bunch of SYN packets, because it keeps all the connection/SN data it gets

- Solution - only add to stack when you've authenticated the client

- Do this by choosing SN a different way

- $SN = F(SRCIP, DSTIP, SRCPORT, DSTPORT, TIMESTAMP, KEY) + \text{timestamp}$

- No need to store the info until after handshake